

Production Environment and the Endpoint

When we discussed the *production environment*, the **endpoint** was defined as the **interface** to the model. This **interface (endpoint)** facilitates an ease of communication between the *model* and the *application*. Specifically, this **interface (endpoint)**

- Allows the *application* to send **user data** to the *model* and
- Receives **predictions** back from the *model* based upon that **user data**.

Model, Application, and Endpoint

```
web_app.py × Application
1 # Main program function defined below that gets user inputs and
2 # returns the model's predictions to the user based upon input data
3 def main():
4     # Retrieve user data
5     input_user_data = get_user_data()
6
7     # Get predictions based upon user's data
8     predictions = ml_model(input_user_data)
9
10    # Displays predictions to user.
11    display_predictions_to_user(predictions)
12
13
14 def ml_model(user_data):
15     """
16     Inputs User Data and returns ML Model's predictions based upon the
17     user input data.
18     Parameters:
19         user_data - the User's input Data that will be used to make
20                     predictions with the model.
21     Returns:
22         models_predictions - the Model's predictions based upon the input
23                             user data.
24     """
25     # Load the user data
26     loaded_data = load_user_data(user_data)
27
```

The image shows a code editor window titled 'web_app.py'. The code defines a `main()` function and an `ml_model(user_data)` function. Annotations with arrows point to specific parts of the code: 'Application' points to the `web_app.py` tab, 'Endpoint' points to the `predictions = ml_model(input_user_data)` line, and 'Model' points to the `def ml_model(user_data):` line.

One way to think of the **endpoint** that acts as this *interface*, is to think of a *Python program* where:

- the **endpoint** itself is like a **function call**
- the **function** itself would be the **model** and

- the *Python program* is the **application**.

The image **above** depicts the association between a *Python program* and the **endpoint**, **model**, and **application**.

- the **endpoint**: *line 8 function call* to ml_model
- the **model**: beginning on *line 14 function definition* for ml_model
- the **application**: *Python program web_app.py*

The image shows a code editor window titled 'web_app.py' with a close button. The code is a Python script for a web application. Annotations with arrows point to specific parts of the code:

- Application**: Points to the file name 'web_app.py'.
- User's Data**: Points to the `input_user_data` argument in the `ml_model` function call.
- Endpoint**: Points to the `predictions = ml_model(input_user_data)` line.
- Model**: Points to the `def ml_model(user_data):` function definition.
- Model's Prediction**: Points to the `predictions` variable on the right side of the assignment.

```
1 # Main program function defined below that gets user inputs and
2 # returns the model's predictions to the user based upon input data
3 def main():
4     # Retrieve user data
5     input_user_data = get_user_data()
6
7     # Get predictions based upon user's data
8     predictions = ml_model(input_user_data)
9
10    # Displays predictions to user.
11    display_predictions_to_user(predictions)
12
13
14 def ml_model(user_data):
15     """
16     Inputs User Data and returns ML Model's predictions based upon the
17     user input data.
18     Parameters:
19     user_data - the User's input Data that will be used to make
20                 predictions with the model.
21     Returns:
22     models_predictions - the Model's predictions based upon the input
23                         user data.
24     """
25     # Load the user data
26     loaded_data = load_user_data(user_data)
27
```

Using this example **above** notice the following:

- Similar to a **function call** the **endpoint** accepts *user data* as the **input** and **returns** the *model's prediction* based upon this **input** through the **endpoint**.
- In the example, the *user data* is the **input argument** and the *prediction* is the **returned value** from the **function call**.

- The **application**, here the *python program*, displays the *model's prediction* to the *application user*.

This example highlights how the **endpoint** itself is just the *interface* between the **model** and the **application**; where this *interface* enables users to get *predictions* from the *deployed model* based on their *user data*.

Next we'll focus on *how* the **endpoint** (*interface*) facilitates communication between **application** and **model**.

Endpoint and REST API

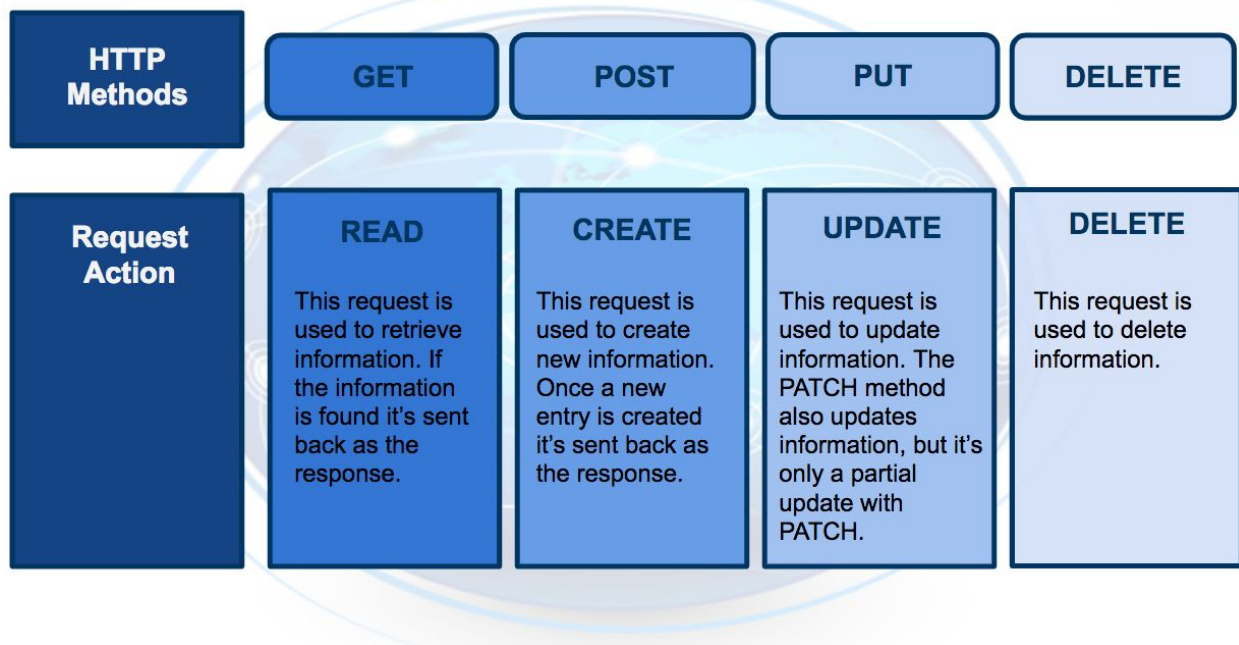
Communication between the **application** and the **model** is done through the **endpoint** (*interface*), where the **endpoint** is an **Application Programming Interface (API)**.

- An easy way to think of an **API**, is as a *set of rules* that *enable* programs, here the **application** and the **model**, to *communicate* with each other.
- In this case, our **API** uses a **RE**presentational **S**tate **T**ransfer, **REST**, architecture that provides a framework for the *set of rules* and *constraints* that must be adhered to for *communication* between programs.
- This **REST API** is one that uses *HTTP requests* and *responses* to enable communication between the **application** and the **model** through the **endpoint** (*interface*).

- Noting that **both** the **HTTP request** and **HTTP response** are *communications* sent between the **application** and **model**.

The **HTTP request** that's sent from your **application** to your **model** is composed of *four* parts:

- **Endpoint**
 - This **endpoint** will be in the form of a URL, Uniform Resource Locator, which is commonly known as a web address.
- **HTTP Method**
 - Below you will find four of the **HTTP methods**, but for purposes of **deployment** our **application** will use the **POST method only**.
- **HTTP Headers**
 - The **headers** will contain additional information, like the format of the data within the message, that's passed to the *receiving* program.
- **Message (Data or Body)**
 - The final part is the **message** (data or body); for **deployment** will contain the *user's data* which is input into the **model**.



The **HTTP response** sent from your model to your application is composed of *three* parts:

- HTTP Status Code
 - If the model successfully received and processed the *user's data* that was sent in the **message**, the status code should start with a **2**, like *200*.
- HTTP Headers
 - The **headers** will contain additional information, like the format of the data within the **message**, that's passed to the receiving program.
- Message (Data or Body)
 - What's returned as the *data* within the **message** is the *prediction* that's provided by the **model**.

This *prediction* is then presented to the *application user* through the **application**. The **endpoint** is the *interface* that *enables communication* between the **application** and the **model** using a **REST API**.

As we learn more about **RESTful API**, realize that it's the **application's** responsibility:

- To format the *user's data* in a way that can be easily put into the **HTTP request message** and *used* by the **model**.
- To translate the *predictions* from the **HTTP response message** in a way that's easy for the *application user's* to understand.

Notice the following regarding the *information* included in the **HTTP messages** sent between **application** and **model**:

- Often *user's data* will need to be in a *CSV* or *JSON* format with a specific *ordering* of the data that's dependent upon the **model** used.
- Often *predictions* will be returned in *CSV* or *JSON* format with a specific *ordering* of the returned *predictions* dependent upon the **model** used.